

# Preventing Gradient Attenuation in Lipschitz Constrained Convolutional Networks

Qiyang Li\*, Saminul Haque\*, Cem Anil, James Lucas,  
Roger Grosse, Jörn-Henrik Jacobsen

January 13, 2020

# Introduction: Adversarial Examples

Given a classifier  $f$  that can correctly classify an input  $\mathbf{x}$  ( $\arg \max f(\mathbf{x}) = t$ ). An adversarial example is a perturbed input that fools the classifier:

$$\arg \max f(\mathbf{x} + \delta \mathbf{x}) \neq t$$

# Introduction: Adversarial Examples

- Image classification (Goodfellow et al., 2014)




$x$

“panda”

57.7% confidence

$+ .007 \times$



$\text{sign}(\nabla_x J(\theta, x, y))$

“nematode”

8.2% confidence

$=$



$x +$

$\epsilon \text{sign}(\nabla_x J(\theta, x, y))$

“gibbon”

99.3 % confidence

# Introduction: Adversarial Examples

- ▶ Image classification (Goodfellow et al., 2014)
- ▶ Text classification (Ebrahimi et al., 2017)

---

South Africa's historic Soweto township marks its 100th birthday on Tuesday in a mood of optimism.  
57% **World**

South Africa's historic Soweto township marks its 100th birthday on Tuesday in a moo**P** of optimism.  
95% **Sci/Tech**

---

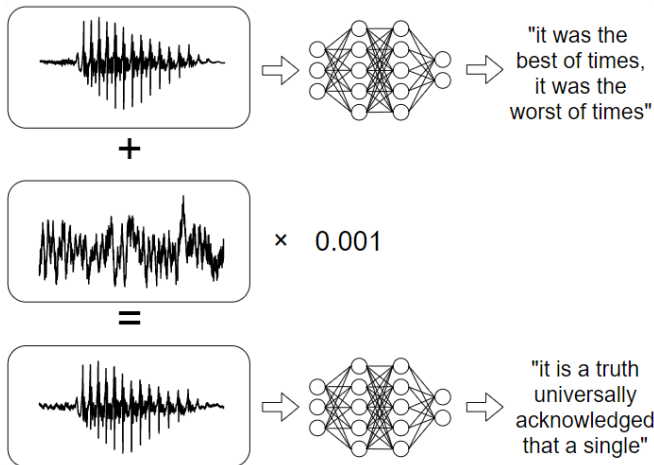
Chancellor Gordon Brown has sought to quell speculation over who should run the Labour Party and turned the attack on the opposition Conservatives.  
75% **World**

Chancellor Gordon Brown has sought to quell speculation over who should run the Labour Party and turned the attack on the o**B**position Conservatives.  
94% **Business**

---

# Introduction: Adversarial Examples

- ▶ Image classification (Goodfellow et al., 2014)
- ▶ Text classification (Ebrahimi et al., 2017)
- ▶ Speech recognition (Qin et al., 2019)



# Introduction: Defense to Adversarial Examples

- ▶ Simple heuristic such as training on adversarial examples can enhance the empirical robustness of the network.

# Introduction: Defense to Adversarial Examples

- ▶ Simple heuristic such as training on adversarial examples can enhance the empirical robustness of the network.
- ▶ Lack of robustness guarantee runs the risk of being attacked in the future ([www.robust-ml.org/defenses](http://www.robust-ml.org/defenses)).

Defense	Venue	Dataset	Threat Model	Natural Accuracy	Claims	Analyses
<a href="#">Mitigating Adversarial Effects Through Randomization</a> (Xie et al.) (code)	ICLR 2018	ImageNet	$\ell_\infty(\epsilon = 10/255)$	99.2% accuracy (on images originally classified correctly by underlying model)	86% accuracy (on images originally classified correctly)	• 0% accuracy [ACW18] (code)
<a href="#">Thermometer Encoding: One Hot Way To Resist Adversarial Examples</a> (Buckman et al.) (code)	ICLR 2018	CIFAR-10	$\ell_\infty(\epsilon = 8/255)$	90% accuracy	79% accuracy	• 30% accuracy [ACW18] (code)
<a href="#">Countering Adversarial Images using Input Transformations</a> (Guo et al.) (code)	ICLR 2018	ImageNet	$\ell_2(\epsilon = 0.06)$	75% accuracy	70% accuracy on ImageNet with average normalized $\ell_2$ perturbation of 0.06	• 0% accuracy [ACW18] (code)
<a href="#">Stochastic Activation Pruning for Robust Adversarial Defense</a> (Dhillon et al.) (code)	ICLR 2018	CIFAR-10	$\ell_\infty(\epsilon = 4/255)$	83% accuracy	51% accuracy	• 0% accuracy [ACW18] (code)

# Introduction: Certifying Adversarial Robustness

- ▶ A function  $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$  is  $K$ -Lipschitz if and only if

$$\|f(\mathbf{x}_1) - f(\mathbf{x}_2)\|_2 \leq K\|\mathbf{x}_1 - \mathbf{x}_2\|_2, \forall \mathbf{x}_1, \mathbf{x}_2 \in \mathbb{R}^n.$$

We denote  $\text{Lip}(f)$  as the smallest value  $K$  such that  $f$  is  $K$ -Lipschitz. This is called the **Lipschitz constant** of  $f$ .

- ▶  $f$  is robust against a  $L_2$  perturbation around an input example  $\mathbf{x}$  within a radius of  $\epsilon$  as long as

$$\sqrt{2}K\epsilon < y_t - \max_{i \neq t} y_i$$

where  $f(\mathbf{x}) = [y_1, y_2, \dots, y_m]$  is the output of the classifier,  $y_t$  is the logit of the correct label.



# Introduction: Training Provably Robust Classifier

- ▶ Provably robust classifier  $f^*$  can be obtained by solving an optimization problem over 1-Lipschitz functions using hinge loss.

$$f^* = \arg \max_{f: \text{Lip}(f) \leq 1} \mathbb{E}_{\mathbf{x} \sim \mathcal{D}} \min \left( 0, y_t - \max_{i \neq t} y_i - \sqrt{2}\epsilon \right)$$

# Introduction: Training Provably Robust Classifier

- ▶ Provably robust classifier  $f^*$  can be obtained by solving an optimization problem over 1-Lipschitz functions using hinge loss.

$$f^* = \arg \max_{f: \text{Lip}(f) \leq 1} \mathbb{E}_{\mathbf{x} \sim \mathcal{D}} \min \left( 0, y_t - \max_{i \neq t} y_i - \sqrt{2}\epsilon \right)$$

- ▶ Need a way to parameterize functions with Lipschitz constraint!

# Introduction: Estimating Wasserstein Distance

- ▶ The 1-Wasserstein distance between two distributions can also be casted as an optimization problem over 1-Lipschitz functions:

$$W_1(P_1, P_2) = \sup_{f: \text{Lip}(f) \leq 1} \left( \mathbb{E}_{\mathbf{x} \sim P_1} [f(\mathbf{x})] - \mathbb{E}_{\mathbf{x} \sim P_2} [f(\mathbf{x})] \right).$$

# Introduction: Estimating Wasserstein Distance

- ▶ The 1-Wasserstein distance between two distributions can also be casted as an optimization problem over 1-Lipschitz functions:

$$W_1(P_1, P_2) = \sup_{f: \text{Lip}(f) \leq 1} \left( \mathbb{E}_{\mathbf{x} \sim P_1} [f(\mathbf{x})] - \mathbb{E}_{\mathbf{x} \sim P_2} [f(\mathbf{x})] \right).$$

- ▶ Need a way to parameterize functions with Lipschitz constraint!

# Background: Parameterizing Functions with Lipschitz Constraint

- ▶ **Idea:** Dynamically scale down  $f$  by  $\text{Lip}(f)$
- ▶ **Problem:** The exact Lipschitz constant of an arbitrarily defined network is extremely challenging to compute.
- ▶ **The Naive Solution:** Use the multiplication of Lipschitz constant for each layer as an upperbound of the global Lipschitz constant.

$$\text{Lip}(f \circ g) \leq \text{Lip}(f) \text{Lip}(g)$$

## Let's pause for a second

- ▶ What if  $\text{Lip}(f \circ g) = \text{Lip}(f) \text{Lip}(g)$  always hold? Can we design the architecture of the networks such that this always holds?

## Let's pause for a second

- ▶ What if  $\text{Lip}(f \circ g) = \text{Lip}(f) \text{Lip}(g)$  always hold? Can we design the architecture of the networks such that this always holds?

## Let's pause for a second

- ▶ What if  $\text{Lip}(f \circ g) = \text{Lip}(f) \text{Lip}(g)$  always hold? Can we design the architecture of the networks such that this always holds?
- ▶ Gradient norm preservation!



# Background: Gradient Norm Preservation

- ▶ A function  $f : \mathbb{R}^n \mapsto \mathbb{R}^m$  is gradient norm preserving if

$$\left\| J^T \mathbf{g} \right\|_2 = \|\mathbf{g}\|_2, \forall \mathbf{g} \in \mathcal{G}.$$

where  $J$  is the input-output Jacobian of  $f$ ,  $\mathcal{G}$  is the set of all possible gradient vectors from the output, which we assume to be  $\mathbb{R}^m$ .


- ▶ A few important properties of GNP functions:
  1. The composition of two GNP functions is GNP.
  2. GNP functions have Lipschitz constant of exactly 1.

## Background: GNP Activation

- ▶ Almost all commonly used activation functions are non-GNP!

# Background: GNP Activation

- ▶ Almost all commonly used activation functions are non-GNP!
- ▶ **Solution:** GroupSort (Anil et al., 2018)



## Background: GNP Linear is Orthogonal

- ▶ If  $f : \mathbb{R}^n \mapsto \mathbb{R}^m$  is linear and GNP, then

$$\mathbf{y} = f(\mathbf{x}) = \mathbf{H}\mathbf{x}$$

with  $\mathbf{H}$  being an matrix with orthonormal rows.

- ▶ Without the loss of generality, we only consider the case where  $\mathbf{H}$  is an orthogonal matrix.
- ▶ Parameterization of/training with orthogonal matrices has been well-studied in the literature.
  1. Björck-Bowie iterative algorithm (Björck and Bowie, 1971).
  2. Matrix exponential of skew symmetric matrix (Lezcano-Casado and Martínez-Rubio, 2019).
  3. Composition of householder matrices (Householder, 1958).
  4. ...

## What about Convolutions?

- ▶ Convolutions are linear. Enforcing GNP on the convolution layer requires the matrix representation of it to be orthogonal.

## What about Convolutions?

- ▶ Convolutions are linear. Enforcing GNP on the convolution layer requires the matrix representation of it to be orthogonal.
- ▶ The orthogonality of a convolutional kernel subject to **cyclic padding** is **invariant** to input size if we assume the input size ( $s$ ) is large relative to the kernel size ( $s \geq 2k - 1$ ).

## What about Convolutions?

- ▶ Convolutions are linear. Enforcing GNP on the convolution layer requires the matrix representation of it to be orthogonal.
- ▶ The orthogonality of a convolutional kernel subject to **cyclic padding** is **invariant** to input size if we assume the input size ( $s$ ) is large relative to the kernel size ( $s \geq 2k - 1$ ).
- ▶ This allows us to study orthogonal convolutions **independent** of the input size.

# What about Convolutions?

- ▶ Convolutions are linear. Enforcing GNP on the convolution layer requires the matrix representation of it to be orthogonal.
- ▶ The orthogonality of a convolutional kernel subject to **cyclic padding** is **invariant** to input size if we assume the input size ( $s$ ) is large relative to the kernel size ( $s \geq 2k - 1$ ).
- ▶ This allows us to study orthogonal convolutions **independent** of the input size.
- ▶ **Definition:** A kernel size  $k$  convolution is orthogonal if

$$\mathcal{T}_{2k-1}^T \mathcal{T}_{2k-1} = \mathbf{I}$$

where  $\mathcal{T}_s$  is the matrix representation of the convolution operation with an input size of  $s$



# Parameterization of Orthogonal Convolution

An orthogonal convolution kernel can be parameterized as a composition of orthogonal convolution kernels with kernel sizes of 1 and 2 (Kautsky and Turcajová, 1994):

$$\mathbf{A} = \mathbf{H} \square [\mathbf{P}_1 \quad (\mathbf{I} - \mathbf{P}_1)] \square \cdots \square [\mathbf{P}_{k-1} \quad (\mathbf{I} - \mathbf{P}_{k-1})]$$

where  $\mathbf{H} \in O(n)$  is an orthogonal matrix of  $n \times n$ , each of  $\mathbf{P}_i$  is a symmetric projector ( $\mathbf{P}_i = \mathbf{P}_i^T = \mathbf{P}_i^2$ ) and  $\square$  represents the composition of convolutions.

## A Surprise from the Topology Standpoint

$$\mathbf{A} = \mathbf{H} \square [\mathbf{P}_1 \quad (\mathbf{I} - \mathbf{P}_1)] \square \cdots \square [\mathbf{P}_{k-1} \quad (\mathbf{I} - \mathbf{P}_{k-1})]$$

- ▶ **Theorem:** The space of 1-D orthogonal convolution kernel has  $2(k-1)n + 2$  (disconnected) **connected components** ( $n$  is the channel size).
- ▶ If we use gradient-based optimization procedure in this space, it can **never** escape the connected component it is initialized in.

## Background: Symmetric Projector

- ▶ Symmetric projectors ( $\mathbf{P} = \mathbf{P}^2 = \mathbf{P}^T$ ) represent orthogonal projections onto linear subspaces.

## Background: Symmetric Projector

- ▶ Symmetric projectors ( $\mathbf{P} = \mathbf{P}^2 = \mathbf{P}^T$ ) represent orthogonal projections onto linear subspaces.
- ▶  $\mathbf{P}_1$  and  $\mathbf{P}_2$  belong to the same connected component iff  $\text{rank}(\mathbf{P}_1) = \text{rank}(\mathbf{P}_2)$ .

## Background: Symmetric Projector

- ▶ Symmetric projectors ( $\mathbf{P} = \mathbf{P}^2 = \mathbf{P}^T$ ) represent orthogonal projections onto linear subspaces.
- ▶  $\mathbf{P}_1$  and  $\mathbf{P}_2$  belong to the same connected component iff  $\text{rank}(\mathbf{P}_1) = \text{rank}(\mathbf{P}_2)$ .
- ▶ The  $n \times n$  symmetric projector space has  $n + 1$  connected components!

Example: Orthogonal Convolution Kernels with  $k = 2$ ,  
 $n = 2$

$$\mathbf{A} = \mathbf{H} \square [\mathbf{P}_1 \quad (\mathbf{I} - \mathbf{P}_1)] \square \cdots \square [\mathbf{P}_{k-1} \quad (\mathbf{I} - \mathbf{P}_{k-1})]$$

Example: Orthogonal Convolution Kernels with  $k = 2$ ,  
 $n = 2$

$$\mathbf{A} = \mathbf{H} \square [\mathbf{P} \quad (\mathbf{I} - \mathbf{P})]$$

Example: Orthogonal Convolution Kernels with  $k = 2$ ,  
 $n = 2$

$$\mathbf{A} = [\mathbf{P} \quad (\mathbf{I} - \mathbf{P})]$$



Example: Orthogonal Convolution Kernels with  $k = 2$ ,  
 $n = 2$

$$\blacktriangleright \mathbf{P} = \begin{bmatrix} x & y \\ y & z \end{bmatrix}.$$

## Example: Orthogonal Convolution Kernels with $k = 2$ , $n = 2$

►  $\mathbf{P} = \begin{bmatrix} x & y \\ y & z \end{bmatrix}.$

►  $\text{rank}(\mathbf{P}) = 0 \rightarrow \mathbf{P} = 0$

## Example: Orthogonal Convolution Kernels with $k = 2$ , $n = 2$

►  $\mathbf{P} = \begin{bmatrix} x & y \\ y & z \end{bmatrix}.$

►  $\text{rank}(\mathbf{P}) = 0 \rightarrow \mathbf{P} = 0$

►  $\text{rank}(\mathbf{P}) = 1 \rightarrow x = 1 - z, y = \pm\sqrt{(1-z)z}, 0 \leq z \leq 1$

## Example: Orthogonal Convolution Kernels with $k = 2$ , $n = 2$

►  $\mathbf{P} = \begin{bmatrix} x & y \\ y & z \end{bmatrix}.$

►  $\text{rank}(\mathbf{P}) = 0 \rightarrow \mathbf{P} = \mathbf{0}$

►  $\text{rank}(\mathbf{P}) = 1 \rightarrow x = 1 - z, y = \pm\sqrt{(1 - z)z}, 0 \leq z \leq 1$

►  $\text{rank}(\mathbf{P}) = 2 \rightarrow \mathbf{P} = \mathbf{I}$


## Example: Orthogonal Convolution Kernels with $k = 2$ , $n = 2$

►  $\mathbf{P} = \begin{bmatrix} x & y \\ y & z \end{bmatrix}$ .

►  $\text{rank}(\mathbf{P}) = 0 \rightarrow \mathbf{P} = \mathbf{0}$

►  $\text{rank}(\mathbf{P}) = 1 \rightarrow x = 1 - z, y = \pm\sqrt{(1-z)z}, 0 \leq z \leq 1$

►  $\text{rank}(\mathbf{P}) = 2 \rightarrow \mathbf{P} = \mathbf{I}$



## A Surprise from the Topology Standpoint

$$\mathbf{A} = \mathbf{H} \square [\mathbf{P}_1 \quad (\mathbf{I} - \mathbf{P}_1)] \square \cdots \square [\mathbf{P}_{k-1} \quad (\mathbf{I} - \mathbf{P}_{k-1})]$$

- ▶ **Theorem:** The space of 1-D orthogonal convolution kernel has  $2(k-1)n + 2$  (disconnected) **connected components** ( $n$  is the channel size).
- ▶ If we use gradient-based optimization procedure in this space, it can **never** escape the connected component it is initialized in.

# Doubling the Channel Size Circumvents the Issue

- **Observation:** For all orthogonal convolutions  $\mathbf{A}^{(n)}$  with channel sizes of  $n$  and arbitrary ranks for  $\mathbf{P}_i^{(n)}$ , there always exists an orthogonal convolution  $\mathbf{A}^{(2n)}$  with a channel size of  $2n$  and  $\text{rank}(\mathbf{P}_i^{(2n)}) = n$  such that

$$\left(\mathbf{A}^{(2n)}\mathbf{x}\right)_{1:n} = \left(\mathbf{A}^{(n)}\mathbf{x}_{1:n}\right), \forall \mathbf{x} \in \mathbb{R}^{2n}$$

# Doubling the Channel Size Circumvents the Issue

- **Observation:** For all orthogonal convolutions  $\mathbf{A}^{(n)}$  with channel sizes of  $n$  and arbitrary ranks for  $\mathbf{P}_i^{(n)}$ , there always exists an orthogonal convolution  $\mathbf{A}^{(2n)}$  with a channel size of  $2n$  and  $\text{rank}(\mathbf{P}_i^{(2n)}) = n$  such that

$$\left(\mathbf{A}^{(2n)}\mathbf{x}\right)_{1:n} = \left(\mathbf{A}^{(n)}\mathbf{x}_{1:n}\right), \forall \mathbf{x} \in \mathbb{R}^{2n}$$

- This means that we can effectively **double** the dimensions in every hidden layer of an neural network to enable it express all connected components of the original network in just **one connected component**!




## Doubling the Channel Size Circumvents the Issue

$$\blacktriangleright n = 1, \mathbf{P}^{(2)} = \begin{bmatrix} x & y \\ y & z \end{bmatrix}.$$

# Doubling the Channel Size Circumvents the Issue


- ▶  $n = 1, \mathbf{P}^{(2)} = \begin{bmatrix} x & y \\ y & z \end{bmatrix}$ .
- ▶ Assume that we parameterize in the connected component where  $\text{rank}(\mathbf{P}^{(2)}) = 1$  (in blue on the left).




$\mathbf{P}^{(2)}$

# Doubling the Channel Size Circumvents the Issue

- ▶  $n = 1$ ,  $\mathbf{P}^{(2)} = \begin{bmatrix} x & y \\ y & z \end{bmatrix}$ .
- ▶ Assume that we parameterize in the connected component where  $\text{rank}(\mathbf{P}^{(2)}) = 1$  (in blue on the left).
- ▶  $x$  can now reach both connected components of  $\mathbf{P}^{(1)}$ : 0 and 1 (as shown on the right).



$\mathbf{P}^{(2)}$



$x$

# Parameterization Summary

- ▶ 1-D orthogonal convolution can be parameterized in the following connected component:

$$\mathbf{A} = \mathbf{H} \square [\mathbf{P}_1 \quad (\mathbf{I} - \mathbf{P}_1)] \square \cdots \square [\mathbf{P}_{k-1} \quad (\mathbf{I} - \mathbf{P}_{k-1})]$$

where  $\mathbf{P}_i \in \mathbb{R}^{2n \times 2n}$  are symmetric projectors with  $\text{rank}(\mathbf{P}_i) = n$ ,  $\mathbf{H} \in O(2n)$ .

- ▶ Similar parameterization also exists in 2-D orthogonal convolutions where the same trick can also be used.

# How does our GNP convolution compare to other convolutions?

- ▶ OSSN: One-sided spectral normalization (Gouk et al., 2018)
- ▶ RKO: Reshaped kernel method (Cisse et al., 2017)
- ▶ SVCM: Singular value clipping and masking (Sedghi et al., 2019)
- ▶ Evaluated on provable adversarial robustness on MNIST and CIFAR-10.

Dataset			OSSN	RKO	SVCM	Ours
MNIST ( $\epsilon = 1.58$ )	Small	Clean	96.86 $\pm$ 0.13	97.28 $\pm$ 0.08	97.24 $\pm$ 0.09	<b>97.54</b> $\pm$ 0.06
		Robust	42.95 $\pm$ 1.09	43.58 $\pm$ 0.44	28.94 $\pm$ 1.58	<b>45.84</b> $\pm$ 0.90
	Large	Clean	98.31 $\pm$ 0.03	98.44 $\pm$ 0.05	97.93 $\pm$ 0.05	<b>98.77</b> $\pm$ 0.05
		Robust	53.77 $\pm$ 1.02	55.18 $\pm$ 0.46	38.00 $\pm$ 1.82	<b>56.66</b> $\pm$ 0.23
CIFAR-10 ( $\epsilon = 36/255$ )	Small	Clean	62.18 $\pm$ 0.66	61.77 $\pm$ 0.63	62.39 $\pm$ 0.46	<b>64.53</b> $\pm$ 0.30
		Robust	48.03 $\pm$ 0.54	47.46 $\pm$ 0.53	47.59 $\pm$ 0.56	<b>50.01</b> $\pm$ 0.21
	Large	Clean	67.51 $\pm$ 0.47	70.01 $\pm$ 0.26	69.65 $\pm$ 0.38	<b>72.41</b> $\pm$ 0.22
		Robust	53.64 $\pm$ 0.49	55.76 $\pm$ 0.16	53.61 $\pm$ 0.51	<b>58.72</b> $\pm$ 0.23

## How does our GNP convolution compare to other convolutions?

- ▶ Estimations of Wasserstein distance between the data and generator distributions of STL-10 GAN and CIFAR-10 GAN.
- ▶ Each estimate is a **strict lower bound**: higher value indicates a tighter estimate.

		OSSN	RKO	Ours
STL-10	MaxMin	$7.39 \pm 0.31$	$8.95 \pm 0.12$	<b><math>9.91 \pm 0.11</math></b>
	ReLU	$7.06 \pm 0.72$	$7.82 \pm 0.21$	<b><math>8.28 \pm 0.19</math></b>
CIFAR-10	MaxMin	$3.29 \pm 0.05$	$4.95 \pm 0.08$	<b><math>5.34 \pm 0.07</math></b>
	ReLU	$3.07 \pm 0.12$	$4.20 \pm 0.06$	<b><math>4.39 \pm 0.07</math></b>

# SOTA Comparison

Dataset		Ours-Large	FC-3	KW-Large (Wong et al., 2018)
<b>MNIST</b> ( $\epsilon = 1.58$ )	Clean	<b>98.77</b> $\pm$ 0.05	98.71 $\pm$ 0.02	88.12
	Robust	<b>56.66</b> $\pm$ 0.23	54.46 $\pm$ 0.30	44.53
<b>CIFAR-10</b> ( $\epsilon = 36/255$ )	Clean	<b>72.41</b> $\pm$ 0.22	62.60 $\pm$ 0.39	59.76
	Robust	<b>58.72</b> $\pm$ 0.23	49.97 $\pm$ 0.35	50.60

- Our GNP convolutional network outperforms the state-of-the-art on certifying deterministic provable adversarial robustness under  $L_2$  metric.

# Take-aways

1. **Gradient norm preserving (GNP)** ConvNets are effective for (1) provably robust image classification, and (2) Wasserstein distance estimation on image domain.
2. Although the orthogonal convolution space is **disconnected**, the issue can be largely mitigated by doubling the channel size and choosing the appropriate connected component to optimize in.
3. We design the **first** GNP ConvNet. It outperforms the state-of-the-art on **deterministic provable adversarial robustness** with  $L_2$  metric on MNIST and CIFAR10.



# References I

- Anil, C., Lucas, J., and Grosse, R. (2018). Sorting out lipschitz function approximation. *arXiv preprint arXiv:1811.05381*.
- Björck, Å. and Bowie, C. (1971). An iterative algorithm for computing the best estimate of an orthogonal matrix. *SIAM Journal on Numerical Analysis*, 8(2):358–364.
- Cisse, M., Bojanowski, P., Grave, E., Dauphin, Y., and Usunier, N. (2017). Parseval networks: Improving robustness to adversarial examples. In *Proceedings of the 34th International Conference on Machine Learning-Volume 70*, pages 854–863. JMLR. org.
- Ebrahimi, J., Rao, A., Lowd, D., and Dou, D. (2017). Hotflip: White-box adversarial examples for text classification. *arXiv preprint arXiv:1712.06751*.
- Goodfellow, I. J., Shlens, J., and Szegedy, C. (2014). Explaining and harnessing adversarial examples. *arXiv preprint arXiv:1412.6572*.

## References II

- Gouk, H., Frank, E., Pfahringer, B., and Cree, M. (2018). Regularisation of neural networks by enforcing Lipschitz continuity. *arXiv preprint arXiv:1804.04368*.
- Householder, A. S. (1958). Unitary triangularization of a nonsymmetric matrix. *Journal of the ACM (JACM)*, 5(4):339–342.
- Kautsky, J. and Turcajová, R. (1994). A matrix approach to discrete wavelets. In *Wavelet Analysis and Its Applications*, volume 5, pages 117–135. Elsevier.
- Lezcano-Casado, M. and Martínez-Rubio, D. (2019). Cheap orthogonal constraints in neural networks: A simple parametrization of the orthogonal and unitary group. *arXiv preprint arXiv:1901.08428*.

## References III

- Qin, Y., Carlini, N., Goodfellow, I., Cottrell, G., and Raffel, C. (2019). Imperceptible, robust, and targeted adversarial examples for automatic speech recognition. *arXiv preprint arXiv:1903.10346*.
- Sedghi, H., Gupta, V., and Long, P. M. (2019). The singular values of convolutional layers. In *International Conference on Learning Representations*.
- Wong, E., Schmidt, F., Metzen, J. H., and Kolter, J. Z. (2018). Scaling provable adversarial defenses. In *Advances in Neural Information Processing Systems*, pages 8400–8409.